
MACHINE LEARNING FOR COMMUNICATION SYSTEMS LAB REPORT MEIC501P

TASK 1: PERFORMANCE ANALYSIS OF SUPERVISED AND UNSUPERVISED LEARNING

NAME: DURGA SITALAKSHMI S

REGISTRATION NO: 24MEC0024

PHONE NUMBER: +91 9731850288

MAIL ID: durgasitalakshmi.s2024@vitstudent.ac.in

TASK 1a: Performance Analysis of Supervised learning

```
#TASK 1(a) : PERFORMANCE ANALYSIS OF SUPERVISED LEARNING
# 24MEC0024
#PDURGA SITALAKSHMI S
#PHONE NUMBER : 9731850288
#MEIC501P

import pandas as pd
# pandas library is imported to do data manipulation and analysis. Making dataframes and changing it as per we want

import numpy as np
# to perform mathematical operations

import librosa
# python library for audio and music analysis. It extracts features from the audio files and process for further use in ML .

from glob import glob
# finds all the pathnames matching a specified pattern

audio_files =glob('Downloads\Actor_01/*.wav')
# collects all .wav extension file from Actor_01 folder from Downloads folder

from matplotlib import pyplot as plt
# Importing pyplot module from matplotlib library for plotting


lst = []
# creates an empty array to append the audio files

length=len(audio_files)
# finds the length of the audio_files

for i in range (length):
    # the commands inside the for loop is to append the audio files to lst array
    y,sr = librosa.load(audio_files[i])
    # The librosa.load() function is used to load an audio file.
    #It returns two values:
    # 1.y, which is a NumPy array containing the audio time series
    # 2.sr, which is the sampling rate of the audio.
    lst.append(y.max()*1000)
    #y.max() computes the highest amplitude in the audio file.
lst.sort()
# sorts the lst in ascending order default. For reverse/descending order , use lst.sort(reverse=True)

df2=pd.DataFrame(lst)
# here the pandas library is called and used to make a dataframe of lst

df2.columns=["freq"]
# naming the column as freq

df2.to_csv("rec.csv")
# converting the dataframe to csv file and saving it as rec.csv file


print(df2)
# printing the dataframe
```

	freq		
0	22.052493	30	70.323378
1	22.440949	31	72.601795
2	26.036292	32	75.142853
3	26.980430	33	81.278026
4	28.652139	34	104.997635
5	29.932763	35	107.536800
6	31.419680	36	107.564509
7	33.317581	37	112.748325
8	35.650913	38	115.064815
9	35.884488	39	116.749868
10	36.971465	40	126.323208
11	38.234800	41	126.476616
12	39.725669	42	130.580217
13	41.008729	43	142.322689
14	41.546766	44	149.627343
15	43.814741	45	166.032284
16	44.969343	46	169.806391
17	45.787793	47	175.272003
18	47.185149	48	178.138345
19	47.394369	49	178.688705
20	50.119177	50	179.265440
21	58.528956	51	202.122718
22	58.665499	52	327.279150
23	58.744207	53	485.229164
24	60.015172	54	506.704867
25	61.159208	55	510.604560
26	61.836861	56	551.321268
27	64.025983	57	579.006553
28	66.383295	58	813.333809
29	67.772545	59	989.103734

Fig 1 : Recorded frequency. Displayed in a dataframe

```

data=pd.read_csv("rec.csv")
#This is a function provided by Pandas to read a CSV (Comma Separated Values) file.
#It returns a DataFrame, which is a 2-dimensional labeled data structure with columns .

Range = data.freq
#assign all frequencies from rec.csv file to Range

y,sr = librosa.load("abnormal.wav")
#Loading the abnormal.wav file

fre = y.max()*1000
#calculating frequency, that is from the all frequency, finding the max amplitude of the signal
# Load the first audio file and its sample rate using librosa
y, sr = librosa.load(audio_files[0])

# Create a plot of the audio time series data
pd.Series(y).plot(figsize=(10, 5), lw=1, title='Test 1 - wave plot')

# Add a grid to the plot (note: this line should call the function with parentheses)
plt.grid()

# Display the plot (note: this line should call the function with parentheses)
plt.show()

# Compute the Short-Time Fourier Transform (STFT) of the audio signal
D = librosa.stft(y)

```

```

# Convert the amplitude of the STFT to decibels
S_db = librosa.amplitude_to_db(np.abs(D), ref=np.max)

# Create a figure and axis for the plot
fig, ax = plt.subplots(figsize=(10, 5))

# Display the spectrogram with a Logarithmic frequency axis
img = librosa.display.specshow(S_db, x_axis='time', y_axis='log', ax=ax)

# Set the title of the spectrogram plot
plt.title("spectrum")

# Display the spectrogram plot
plt.show()

print(" The given audiosfrequency is ",fre)
# printing the frequency of the uploaded 'abnormal.wav' file

if(fre>0 and fre<Range[30]):
    print("THE GIVEN SAMPLE IS CALM")
# if the fre is greater than 0 and less than 30 ,then the audio signal is Calm.
# here all the frequencies that we saved in rec.csv file is saved in Range, so comparing the value from the range of frequencies.

elif(fre>Range[30] and fre<Range[50]):
    print("THE GIVEN SAMPLE IS NORMAL")
# if the fre is greater than 30 and Less than 50 ,then the audio signal is Normal.

elif(fre>Range[50]):
    print("THE GIVEN SAMPLE IS ABNORMAL")
# if the fre is greater than 50 then the audio signal is Abnormal.

```

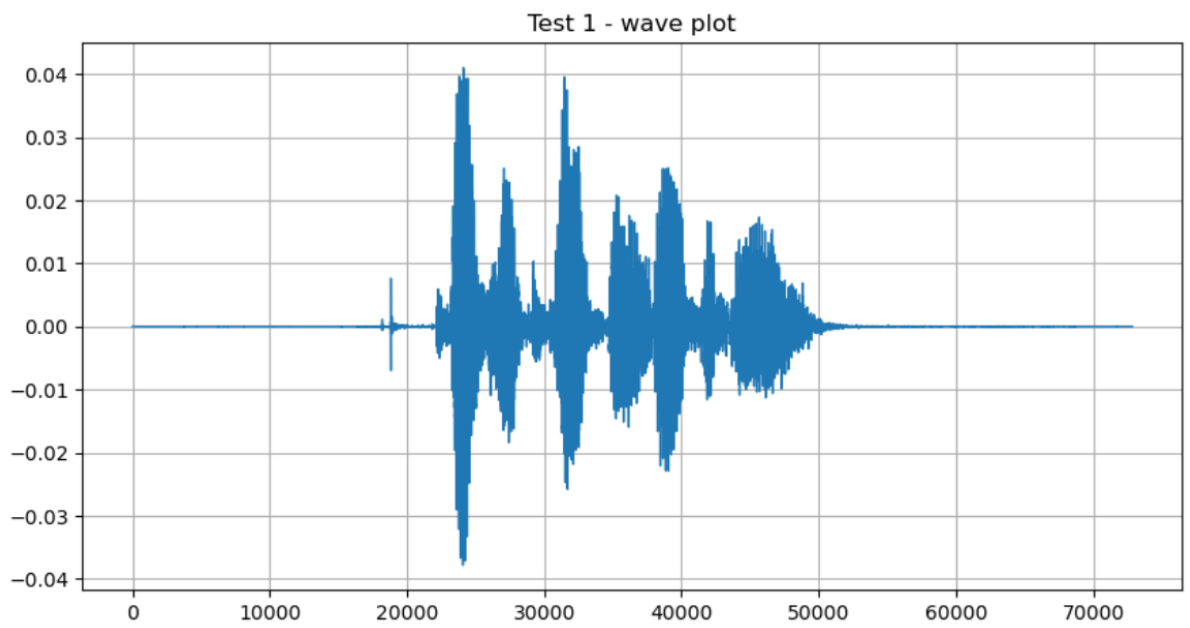


Fig 2 : The waveform of the signal(abnormal wave considered)




Fig 3 : The spectrum of the signal or waveform

The given audiosfrequency is 241.27447605133057
THE GIVEN SAMPLE IS ABNORMAL

Fig 5 : The output showing the frequency and the type of audio signal based on the given conditions

INFERENCE

Feature Extraction:

The code uses `librosa.load` to get the maximum amplitude from audio files.

Data Preparation:

The extracted amplitudes are stored, sorted, converted into a DataFrame, and saved as a CSV file.

Labelling:

Audio samples are labelled as "Calm," "Normal," or "Abnormal" based on amplitude ranges, creating target variables for supervised learning.

Classification:

The provided audio sample (`abnormal.wav`) is classified by comparing its maximum amplitude to the stored ranges, simulating how a supervised model makes predictions.

TASK 1b: Performance Analysis of Unsupervised learning – K means

```
#TASK 1(b) : PERFORMANCE ANALYSIS OF UNSUPERVISED LEARNING
# 24MEC0024
#DURGA SITALAKSHMI S
#PHONE NUMBER : 9731850288
#MEIC501P

import pandas as pd
# Importing the pandas library for data manipulation and analysis

import librosa
# Importing the librosa library for audio and music processing

import numpy as np
# Importing the numpy library for numerical operations on arrays

from sklearn.cluster import KMeans
# Importing the KMeans clustering algorithm from scikit-learn library

from sklearn.preprocessing import StandardScaler
# Importing StandardScaler for feature scaling from scikit-learn library

from matplotlib import pyplot as plt
# Importing pyplot module from matplotlib library for plotting

from glob import glob
# Importing glob module to find all pathnames matching a specified pattern

import warnings
# Importing warnings library to manage warnings
```

```
warnings.filterwarnings("ignore")
# Ignores all warnings to prevent them from being displayed

audio_files = glob('Downloads\Audio_files/*.wav')
# Finds all .wav audio files in the specified directory using glob

mfccs = []
# Initializes an empty list to store MFCC (Mel-frequency cepstral coefficients) features

spectral_centroid = []
# Initializes an empty list to store spectral centroid features

l = len(audio_files)
# Stores the number of audio files found in the directory

print(l)
```

<>:34: SyntaxWarning: invalid escape sequence '\A'
<>:34: SyntaxWarning: invalid escape sequence '\A'
C:\Users\sduhg\AppData\Local\Temp\ipykernel_19060\2430285894.py:34: SyntaxWarning: invalid escape sequence '\A'
audio_files = glob('Downloads\Audio_files/*.wav')

6

```
# Initialize an empty list for MFCCs and spectral centroids
mfccs = []
spectral_centroid = []

# Iterate over the range of audio files
for i in range(l):
    # Load the audio file and its sample rate
    animal, sr = librosa.load(audio_files[i])
    # Extract MFCC features with 20 coefficients
    mfccs_anim = librosa.feature.mfcc(y=animal, sr=sr, n_mfcc=20)
    # Extract spectral centroid features
    spectral_centroids_anims = librosa.feature.spectral_centroid(y=animal, sr=sr)
    # Append the maximum value of the 10th MFCC coefficient to the list
    mfccs.append(max(mfccs_anim[9]))
    # Append the maximum spectral centroid value to the list
    spectral_centroid.append(max(max(spectral_centroids_anims)))

# Create a DataFrame with MFCC and spectral centroid values
df = pd.DataFrame()
df['mfcc'] = mfccs
df['spectral'] = spectral_centroid

# Generate descriptive statistics for the DataFrame
stats_summary = df.describe()

# Standardize the MFCC and spectral centroid values
scaler = StandardScaler()
df[['mfcc_t', 'spectral_t']] = scaler.fit_transform(df[['mfcc', 'spectral']])
```



```

# Apply KMeans clustering with 2 clusters
KM = KMeans(n_clusters=2)
y_predict = KM.fit_predict(df[['mfcc', 'spectral']])

# Add the cluster labels to the DataFrame
df['cluster'] = y_predict

# Print the DataFrame with the cluster labels
print(df)

# Plot the standardized MFCC vs. spectral centroid with cluster coloring
plt.scatter(df['mfcc_t'], df['spectral_t'], c=df['cluster'])
#This line plots the standardized MFCC values (df['mfcc_t']) against the standardized spectral centroid values (df['spectral_t']).
#The c=df['cluster'] argument uses the cluster labels generated by the KMeans algorithm to color the points.
#Each unique value in df['cluster'] will be assigned a different color.
plt.grid()
plt.xlabel('MFCC')
plt.ylabel('SPECTRAL_CENTROID')
plt.title("CLUSTER PLOT")
plt.show()

# Generate descriptive statistics again (if needed for later use)
stats_summary = df.describe()

```

```

# Create a dictionary to store the summary statistics
summary_stats = {
    'mean': stats_summary.loc['mean', ['mfcc', 'spectral', 'mfcc_t', 'spectral_t', 'cluster']],
    'std': stats_summary.loc['std', ['mfcc', 'spectral', 'mfcc_t', 'spectral_t', 'cluster']],
    'min': stats_summary.loc['min', ['mfcc', 'spectral', 'mfcc_t', 'spectral_t', 'cluster']],
    '25%': stats_summary.loc['25%', ['mfcc', 'spectral', 'mfcc_t', 'spectral_t', 'cluster']],
    '50%': stats_summary.loc['50%', ['mfcc', 'spectral', 'mfcc_t', 'spectral_t', 'cluster']],
    '75%': stats_summary.loc['75%', ['mfcc', 'spectral', 'mfcc_t', 'spectral_t', 'cluster']],
    'max': stats_summary.loc['max', ['mfcc', 'spectral', 'mfcc_t', 'spectral_t', 'cluster']]
}

# Convert the dictionary to a DataFrame
summary_df = pd.DataFrame(summary_stats)

# Print the summary DataFrame
print(summary_df)

```

	mfcc	spectral	mfcc_t	spectral_t	cluster
0	68.688110	2410.670365	2.071106	-0.496861	1
1	33.581944	2185.434301	0.134933	-0.907513	1
2	31.278986	1973.660512	0.007921	-1.293619	1
3	19.093605	3562.220285	-0.664126	1.602649	0
4	19.957207	3098.451969	-0.616497	0.757105	0
5	14.212361	2868.710377	-0.933336	0.338239	0

Fig 6 : The dataframe of the mfcc, mfcc_t , spectrum , spectrum_t and the cluster category

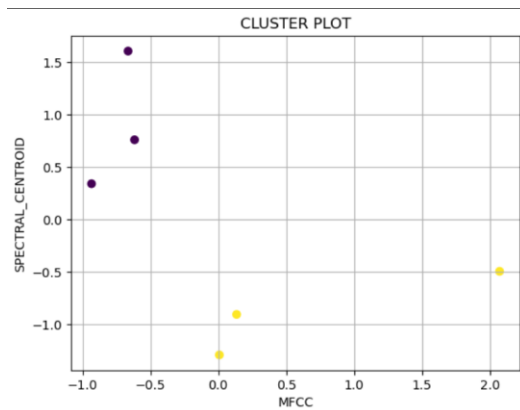


Fig 7 : The cluster plot

	mean	std	min	25%	50% \
mfcc	3.113537e+01	19.862318	14.212361	19.309505	25.618096
spectral	2.683191e+03	600.835141	1973.660512	2241.743317	2639.690371
mfcc_t	-3.700743e-17	1.095445	-0.933336	-0.652219	-0.304288
spectral_t	-4.163336e-16	1.095445	-1.293619	-0.804850	-0.079311
cluster	5.000000e-01	0.547723	0.000000	0.000000	0.500000

	75%	max
mfcc	33.006204	68.688110
spectral	3041.016571	3562.220285
mfcc_t	0.103180	2.071106
spectral_t	0.652389	1.602649
cluster	1.000000	1.000000

Fig 8 : The mean, standard deviation of mfcc, spectral, Spectral_t and cluster

```
# Load the first audio file and its sample rate using librosa
y, sr = librosa.load(audio_files[0])

# Create a plot of the audio time series data
pd.Series(y).plot(figsize=(10, 5), lw=1, title='Test 1 - wave plot')

# Add a grid to the plot (note: this line should call the function with parentheses)
plt.grid()

# Display the plot (note: this line should call the function with parentheses)
plt.show()

# Compute the Short-Time Fourier Transform (STFT) of the audio signal
D = librosa.stft(y)

# Convert the amplitude of the STFT to decibels
S_db = librosa.amplitude_to_db(np.abs(D), ref=np.max)

# Create a figure and axis for the plot
fig, ax = plt.subplots(figsize=(10, 5))

# Display the spectrogram with a Logarithmic frequency axis
img = librosa.display.specshow(S_db, x_axis='time', y_axis='log', ax=ax)

# Set the title of the spectrogram plot
plt.title("spectrum")

# Display the spectrogram plot
plt.show()
```




Fig 9 : Wave Plot




Fig 10 : Spectrum Plot

```

# Compute MFCC (Mel-frequency cepstral coefficients) from the audio signal
mfc = librosa.feature.mfcc(y=y, sr=sr)

# Create a new figure for the plot with specified figure size
plt.figure(figsize=(10, 4))

# Display the MFCC spectrogram with time on the x-axis
librosa.display.specshow(mfc, x_axis='time')

# Add a colorbar to the plot with formatting for the decibel scale
plt.colorbar(format='%+2.0f dB')

# Set the title of the plot to 'MFCCs'
plt.title('MFCCs')

# Display the plot
plt.show()

```




Fig 11 : MFCC Spectrum

INFERENCE

Feature Extraction:

The code uses librosa to extract MFCC and spectral centroid features from audio files. MFCCs capture essential spectral characteristics, while spectral centroids indicate the spectrum's "center of mass."

Data Preparation:

Extracted features are stored in lists, converted to a DataFrame, and standardized using StandardScaler, ensuring comparable feature scales for better clustering performance.

Clustering:

KMeans clustering groups the standardized features into two clusters, uncovering inherent patterns in the audio data without predefined labels.

Visualization:

The clustered data is visualized with a scatter plot of MFCC and spectral centroid features, illustrating cluster separation and groupings.

Spectrogram and MFCC Visualization:

The code plots the waveform, spectrogram (via STFT), and MFCCs of a sample audio file, providing detailed insights into the audio signal's structure for further analysis and cluster interpretation.

Task 1c : Supervised Learning - Analysing Support Vector Machine working

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import os
import shutil
import matplotlib.pyplot as plt
import numpy as np
import librosa
import librosa.display
import IPython.display as ipd

# Output directory to clear
output_dir = "op"

# Clear the contents of the output directory
shutil.rmtree(output_dir, ignore_errors=True)
os.makedirs(output_dir, exist_ok=True)

print(f"Contents of {output_dir} cleared.")

import librosa
import soundfile as sf

# Path to the dataset
dataset_path = "Downloads/16000_pcm_speeches"
```

```

# Output directory to save the combined files
output_dir = "op"

# Create the output directory if it doesn't exist
os.makedirs(output_dir, exist_ok=True)

# List of speaker folders
speaker_folders = [
    "Benjamin_Netanyau",
    "Jens_Stoltenberg",
    "Julia_Gillard",
    "Magaret_Tarcher",
    "Nelson_Mandela"
]

# Number of files to combine for each speaker
num_files_to_combine = 120

# Iterate over each speaker's folder
for speaker_folder in speaker_folders:
    speaker_folder_path = os.path.join(dataset_path, speaker_folder)

    # List the first num_files_to_combine WAV files in the speaker's folder
    wav_files = [f"{i}.wav" for i in range(num_files_to_combine)]

    # Combine all WAV files into a single long file
    combined_audio = []
    for wav_file in wav_files:
        wav_file_path = os.path.join(speaker_folder_path, wav_file)
        audio, sr = librosa.load(wav_file_path, sr=None)
        combined_audio.extend(audio)

    # Save the combined audio file
    output_file_path = os.path.join(output_dir, f"{speaker_folder}_combined.wav")
    sf.write(output_file_path, combined_audio, sr)

print("Combination complete. Combined files saved in:", output_dir)

```

```

Contents of op cleared.
Combination complete. Combined files saved in: op

```

Fig 12 :Output of the first part of the code to combine the files

```

import tensorflow as tf
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler

# Set the parent directory for speaker folders
parent_dir = "Downloads/16000_pcm_speeches"

# List of speaker folders
speaker_folders = [
    "Benjamin_Netanyau",
    "Jens_Stoltenberg",
    "Julia_Gillard",
    "Magaret_Tarcher",
    "Nelson_Mandela"
]

```

```

def extract_features(parent_dir, speaker_folders):
    features = []
    labels = []

    for i, speaker_folder in enumerate(speaker_folders):
        speaker_folder_path = os.path.join(parent_dir, speaker_folder)

        for filename in os.listdir(speaker_folder_path):
            if filename.endswith(".wav"):
                file_path = os.path.join(speaker_folder_path, filename)
                audio, sr = librosa.load(file_path, sr=None, duration=1)
                mfccs = librosa.feature.mfcc(y=audio, sr=sr, n_mfcc=13)

                # Normalize MFCC features
                mfccs = StandardScaler().fit_transform(mfccs)
                features.append(mfccs.T)
                labels.append(i)

    return np.array(features), np.array(labels)

```



```

# Extract features and Labels
X, y = extract_features(parent_dir, speaker_folders)

# Print the first few features
for feature in X[:1]:
    print(feature)

from tensorflow.keras.callbacks import EarlyStopping
from sklearn.preprocessing import LabelEncoder

# Encode labels with explicit classes
label_encoder = LabelEncoder()
y = label_encoder.fit_transform(y)
label_encoder.classes_ = np.array(speaker_folders)

# Split the data into training, validation, and test sets
#X_train, X_temp, y_train, y_temp = train_test_split(X, y, test_size=0.2, random_state=42)
#X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=42)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

```

```

# Print the shapes of training and validation data
print("Training Data Shape:", X_train.shape)
print("Test Data Shape:", X_test.shape)
#print("Validation Data Shape:", X_val.shape)

```

```

[[-3.46410155e+00  2.88717210e-01  2.88714826e-01  2.88710833e-01
  2.88705289e-01  2.88698137e-01  2.88689405e-01  2.88679123e-01
  2.88667232e-01  2.88653851e-01  2.88638890e-01  2.88622409e-01
  2.88604409e-01]
 [-3.46410155e+00  2.88697690e-01  2.88696378e-01  2.88694263e-01
  2.88691312e-01  2.88687497e-01  2.88682818e-01  2.88677305e-01
  2.88670957e-01  2.88663775e-01  2.88655698e-01  2.88646817e-01
  2.88637102e-01]
 [-3.46410179e+00  2.88694620e-01  2.88693517e-01  2.88691700e-01
  2.88689107e-01  2.88685828e-01  2.88681775e-01  2.88677037e-01
  2.88671494e-01  2.88665295e-01  2.88658351e-01  2.88650692e-01
  2.88642257e-01]
 [-3.41496515e+00  4.01689023e-01 -5.24658822e-02  6.99484289e-01
  2.78740406e-01  2.13533744e-01  3.30969393e-01  3.79182965e-01
  6.90628961e-02  2.89402425e-01  2.54391849e-01  3.14149350e-01
  2.36824691e-01]
 [-3.32731915e+00  5.34021258e-01 -3.14830184e-01  9.45635557e-01
  3.05575788e-01  2.24249974e-01  2.92698741e-01  3.76336753e-01

```

Fig 13 :Output of the features

```
Training Data Shape: (6000, 32, 13)
Test Data Shape: (1501, 32, 13)
```

Fig 14 :The output of the shape of the training and testing set

```
from sklearn.mixture import GaussianMixture
from sklearn.svm import SVC
from sklearn.model_selection import train_test_split

# Split the data into training and test sets
#X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train Gaussian Mixture Model (GMM)
#gmm = GaussianMixture(n_components=10, covariance_type='diag')
#gmm.fit(X_train, y_train)

# Extract features using GMM
#X_train_features = gmm.predict_proba(X_train)
#X_test_features = gmm.predict_proba(X_test)
```

```
from sklearn.mixture import GaussianMixture
from sklearn.svm import SVC
from sklearn.model_selection import train_test_split

# Split the data into training and test sets
#X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train Gaussian Mixture Model (GMM)
#gmm = GaussianMixture(n_components=10, covariance_type='diag')
#gmm.fit(X_train, y_train)

# Extract features using GMM
#X_train_features = gmm.predict_proba(X_train)
#X_test_features = gmm.predict_proba(X_test)
```

```

import numpy as np

# Flatten the input data
X_train_flat = X_train.reshape(X_train.shape[0], -1)
X_test_flat = X_test.reshape(X_test.shape[0], -1)

# Initialize and train the SVM classifier
svm_classifier = SVC(kernel='rbf', C=1.0, gamma='scale')
svm_classifier.fit(X_train_flat, y_train)

from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Make predictions on the test set
y_pred = svm_classifier.predict(X_test_flat)

```

```

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print("Test Accuracy:", accuracy)

# Compute confusion matrix
confusion_mat = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(confusion_mat)

# Generate classification report
class_report = classification_report(y_test, y_pred)
print("Classification Report:")
print(class_report)

```

```

import seaborn as sns

# Plot the confusion matrix
plt.figure(figsize=(7, 5))
sns.heatmap(confusion_mat, annot=True, fmt="d", cmap="Blues", xticklabels=speaker_folders, yticklabels=speaker_folders)

# Rotate x-axis labels by 45 degrees
plt.xticks(rotation=45, ha="right")

plt.title("Confusion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()

```

Test Accuracy: 0.9713524317121919

Confusion Matrix:

```

[[304  4  0  2  0]
 [ 24 283  2  1  0]
 [  2  0 280  1  0]
 [  3  3  1 276  0]
 [  0  0  0  0 315]]

```

Classification Report:

	precision	recall	f1-score	support
0	0.91	0.98	0.95	310
1	0.98	0.91	0.94	310
2	0.99	0.99	0.99	283
3	0.99	0.98	0.98	283
4	1.00	1.00	1.00	315
accuracy			0.97	1501
macro avg	0.97	0.97	0.97	1501
weighted avg	0.97	0.97	0.97	1501

Fig 15 : Output of the test accuracy, confusion matrix, classification report and average




Fig 16 :Output showing the graph of the confusion matrix – heat map

INFERENCE

Feature Extraction:

This code uses the librosa library to extract features from audio files, specifically focusing on MFCC (Mel-Frequency Cepstral Coefficients) features. These MFCCs are crucial for capturing the unique characteristics of the audio signals.

Data Preparation:

The extracted MFCC features are collected in lists and then converted into a NumPy array. This structured data is split into training and testing sets using `train_test_split`, ensuring that the model is trained and evaluated on different subsets of data.

Label Encoding:

The labels representing different speakers are transformed into numerical format using Label Encoder. This conversion is necessary for the machine learning model to process and understand the categorical data.

Data Flattening:

The input data is reshaped into a 2D array where each row represents a sample with its corresponding features. This format is required by the SVM classifier.

SVM Classification:

An SVM (Support Vector Machine) classifier is initialized and trained using the training data. The training process involves finding the optimal hyperplane that separates the classes (different speakers) in the feature space.

Prediction and Evaluation:

The trained SVM model makes predictions on the test set. The model's performance is evaluated using various metrics:

- **Accuracy Score:** Measures the proportion of correctly classified samples.
- **Classification Report:** Provides detailed metrics such as precision, recall, and F1-score for each class.
- **Confusion Matrix:** Shows the number of correct and incorrect predictions for each class, offering deeper insights into the model's performance.

Visualization:

A heatmap of the confusion matrix is created to visualize the classifier's performance. The heatmap illustrates the true labels on the y-axis and the predicted labels on the x-axis, with color intensity indicating the number of samples in each category.

